

Introduction

Background and Motivation

- Parallel computing is a part of HPC.
 - HPC also includes everything else that makes the computation fast.
 - No point parallelizing without increasing performance.
 - You might want to optimize for the architecture.
 - Sometimes overhead outweighs benefits from parallelization.
- Focusing on parallel algorithms.
 - Different version of parallel algorithms suits different architecture or models.
- Many application yo.
- People made super computers throughout the 1900s
- Super computers rely on carefully designed interconnects.
- Cloud computers are just AWS instances.
- Many aspects

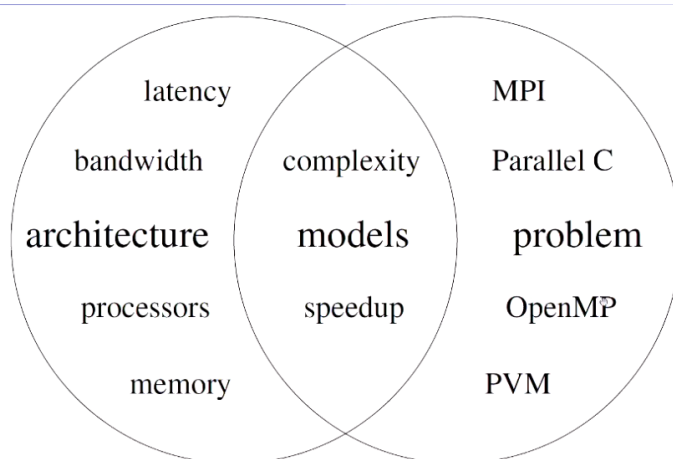


Figure: Overlapping aspects of parallel computing.

Complexity

- $f(n) = O(g(n)) \Rightarrow f$ grows no faster than g
- $f(n) = \Omega(g(n)) \Rightarrow f$ grows no slower than g
- $f(n) = o(g(n)) \Rightarrow f$ grows slower than g
- $f(n) = \omega(g(n)) \Rightarrow f$ grows faster than g
- $f(n) = \Omega(g(n)) \wedge f(n) = O(g(n)) \Rightarrow f(n) = \Theta(g(n))$
- Strictly speaking we should really use \in instead of $=$
- Some common name for complexities:
 - Constant
 - Logarithmic
 - Polylog: $(\log(n))^c$
 - Linearithmic: $n \log n$
 - Quadratic: n^2
 - Polynomial or geometric
 - Exponential
 - Factorial
- Log factor are often ignored.

Model

- RAM model: *random access machine*
 - Common model when we talk about sequential time complexity.
- Multiplying the number of computers by a constant factor doesn't change the complexity.
 - Solution: allow p , the number of processors to increase with problem size and hence reduces the complexity.

PRAM

- Parallel Random Access Machine
- p number of RAM processors, each have private memory and share a large shared memory, all memory access takes the same amount of time.
- Does things synchronously, AKA in lock steps.
- PRAM pseudo code looks like regular pseudo code but there's this
for $i \leftarrow 0$ **to** $n - 1$ **do in parallel**
 processor i **does** thingy

Many different PRAM model

- EREW: exclusive read, exclusive write
- CREW: concurrent read, exclusive write
- CRCW: concurrent read, concurrent write
 - Concurrent write have different types
 - COMMON: Error when two processor tries to write to the same location with different value.
 - ARBITRARY: Pick a arbitrary processor if many processor writes the same time.
 - PRIORITY: Processor with lowest ID writes.
 - COMBINING: Runs a function whenever multiple processors tries to write at the same time.
 - Too powerful.
- ERCW: exclusive read, concurrent write (never used)

Power of model: expresses the set of all problems that can be solved within a certain complexity.

- A is more powerful than B if A can solve a larger set of problems within any complexities.
- A is equally powerful as B if they can solve the same set problems within any complexities.
- Partial ordering.
- COMMON, ARBITRARY, PRIORITY and COMBINING are in increasing order of power.
- Any CRCW PRIORITY PRAM can be simulated by a EREW PRAM with a complexity increase of $\mathcal{O}(\log p)$
- *Parallel Computation Thesis*: any thing can be solved with a Turing Machine with polynomially bounded space can be solved in polynomially bounded space with unlimited processors.
 - Unbounded *word sizes* are not useful, so we limit word counts to $\mathcal{O}(\log p)$
- *Nick's Class* (NC): Solvable in polylog time with poly number of processors.
- Widely believed that $\mathbf{NP} \neq \mathbf{P}$

Definitions (need to remember)

- $w(n) = t(n) \times p(n)$ where $w(n)$ is the work / cost, $t(n)$ is the time and $p(n)$ is the number of processors.
- Optimal processor allocation means: $t(n) \times p(n) = \Theta(T(n))$ where $T(n)$ is the time taking by a sequential algorithm.
 - Equivalent to $t(n) \times p(n) = O(T(n))$ because $t(n) \times p(n) = \Omega(T(n))$ always.
- $\text{Speedup}(n) = \frac{T(n)}{t(n)}$
 - Speedup optimal = processor optimal.
- Optimal: processor optimal AND $t(n) = \mathcal{O}(\log^k n)$
 - Processor optimal and polylog in time.
- Efficient: Assume $T(n) = \Omega(n)$ $w(n) = \mathcal{O}(T(n) \log^\alpha n)$ AND polylog in time
 - Optimal but polylog increase in work.
- **size**: $\text{Size}(n)$ is the total number of operations it does.
- **efficiency**: $\eta(n)$ speedup per processor
 - $\eta(n) = \frac{T(n)}{w(n)} = \frac{\text{Speedup}(n)}{p(n)}$
- You can decrease p and increase t by a factor of $\mathcal{O}\left(\frac{p_1}{p_2}\right)$, $w(n)$ doesn't increase its complexity.
 - Can't do it the other way around.

Brent's Principle (important)

- If something can be done with size x and t time with infinite processors, then it can be done in $t + \frac{x-t}{p}$ time with p processors

Amdahl's Law

- Maximum speedup: if f is the fraction of time that can't be parallelized, then $\text{Speedup}(p) \rightarrow \frac{1}{f}$ as $p \rightarrow \infty$
 - Honestly very obvious.

Gustafson's Law

- s is fraction time of serial part, r is fraction time of parallel part, then $\text{Speedup}(p) = \Omega(p)$
 - Very obvious again...

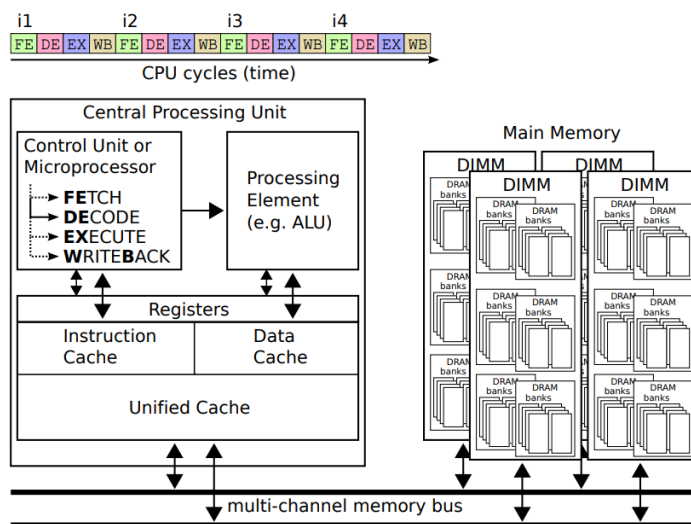
Algorithms

- sum
- logical or
- Maximum
 - n^2 processors all compare all elements and set is_max array to false if element isn't maximum.
 - Only processor with element being max write it to the returning memory address.
- Maximum n^2
 - $\mathcal{O}(\log \log n)$
 - n processor on n elements.
 - Is efficient
 - Make elements into a square, find maximum on each row recursively.
 - Find maximum of maximum of the rows using maximum.
 - $\mathcal{O}(\log \log n)$ levels of recursion, each level takes $\mathcal{O}(1)$ times
- Element Uniqueness
 - Have an array size of MAX_INT.
 - Write processor ID to the array with the element.
 - Check if processor ID is indeed there, if not there's another element there.
- Replication

- $O(\log n)$
- Replication optimal
 - $p = \frac{n}{\log(n)}$ and copy at the end.
- Broadcast
 - Just replicate
- Simulate PRIORITY with COMMON n^2
 - Minimum version of Maximum
- Simulate PRIORITY with EREW
 - All processor wants to write
 - Sort array A of tuples (address, processorID) using Cole's Merge Sort.
 - For each processor k, if $A[k].\text{address} \neq A[k-1].\text{address}$ then $A[k].\text{processorID}$ is the smallest ID that wants to write to that address.

Architecture

- Fetch Decode Execute WriteBack



- Bus is a wire and everyone can see everything on that wire.
- Pipeline: let's do all of them at the same time for the next 4 instructions
 - Need to predict the next 4 instructions sometimes.
- Superpipeline: Do all of them for the next 8 (or more) instructions.
- Superscalar: Multiple pipeline in parallel
- Word size: 64 bits, 32 bits etc, various aspects:
 - Integer size
 - Float size
 - Instruction size
 - Address resolution (mostly bytes)
- Single instruction multiple data SIMD
 - Make word size more complicated
- Coprocessor
 - Used to means stuff directly connected to the CPU like a floating point processor.
 - Now can means FPGA or GPU.
- Multicore processor are just single core duplicated but they all have one extra single shared cache.

- Classification of parallel architectures
 - SISD regular single core.
 - SIMD regular modern single core.
 - MIMD regular multicore.
 - MISD doesn't exist.
- SIMD vs MIMD
 - Effectively SIMD vs non-SIMD
 - Most processor have multicore and SIMD on each core.
 - So a balance between the two.
 - SIMD cores are larger so less of them fit on a die.
 - SIMD is faster at vector operations.
 - SIMD is not useful all the time so sometimes the SIMD part sit idle.
 - SIMD is harder to program.
- Shared memory: All memory can be accessed by all processors.
 - All memory access truly equal time: symmetric multi-processor.
 - Only can have so many cores when the bus is only so fast.
 - Making more buses doesn't help cause space also slows things down.
 - Sometimes can be done with switching interconnect network.
 - Some processor access some memory faster.
 - More complex network.
 - Distributed shared memory: each processor have its own memory but interconnect network exist so you can read other people's memory.
 - *non-uniform memory access* NUMA
 - Static interconnect network: each node connect to some neighbors.
 - *degree*: just like degree in graphs.
 - *diameter*: just like in graphs.
 - *cost* = degree × diameter
- Distributed memory: Each processor have its own memory. Each process live on one processor.
- Blade contains Processor / Package / Socket which contains Core which contains ALU.
- Implicit vs explicit: explicit → decision made by programmer
 - Parallelism: Can I write a sequential algorithm.
 - Decomposition: Can I pretend threads processes doesn't exist.
 - Mapping: Can I pretend all cores are the same.
 - Communication.

- Single Program Multiple Data: one exe
- Multiple Program Multiple Data: multiple exe

Other HPC considerations

- Cache friendliness
- Processor-specific code
- Compiler optimization.
 - Compiler from CPU maker are usually better.
 - So Intel compiler is better than both clang and gcc.

Memory interleaving

- Memory module takes a while to recharge, so we interleave a page on different memory module.